

CS231n

Lecture 1,2 summary

History of computer vision study

Effort to recognize object

- Simplify
 - Block World(Larry Roberts,1963)
 - Stages of Visual Representation(David Marr,1970s)
 - Generalized Cylinder(Brooks&Binford, 1979)
 - Pictorial Structure(Fischler and Elschlager, 1973)
 - Constructing Line and Edges(David Lowe, 1987)
- object recognition was difficult. New approach image seg.
- Normalize cut (Shi & Malik, 1997)


History of computer vision study

Effort to recognize object

- Feature based recognition was tried.
- Machine Learning Algorithm(SVM, Boosting, CNN,Etc) appeared
→ face detection..
- “SIFT” & Object Recognition (David Lowe, 1999)
- Spatial Pyramid Matching(Lazebnik, Schmid & Ponce, 2006)
- Histogram of Gradients (HoG) (Dalal & Triggs, 2005)
- Deformable Part Model (Felzenswalb, McAllester, Ramanan, 2009)

Some changes

- Quality of Pictures
 - Computing power
 - Large Data
- CNN have become an important tool for object recognition



(Lec 2) Image Classification

- Core task in computer vision
- Find fit label in label sets.
- Easy to human, hard to machine
- Semantic Gap

Challenges

- Viewpoint variation
 - Illumination
 - Deformation
 - Occlusion
 - Background Clutter
 - Intraclass variation
- NO obvious algorithm

Data Driven Approach


- 1. Collect a dataset of images and labels
- 2. Use Machine Learning to train a classifier
- 3. Evaluate the classifier on new images

```
def train(images, labels):  
    # Machine learning!  
    return model
```


```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```

First Classifier : Nearest Neighbor


- Metric distance : L1,L2
- Train : $O(1)$ Predict $O(N)$
- K-Nearest Neighbor : Take majority vote from K closest points
- Hyperparameter : K value, Metric define.




K = 1



K = 1



K = 3




K = 5

Setting hypermeter

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data




Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!



Idea #4: Cross-Validation: Split data into **folds**, try each fold as validation and average the results



Useful for small datasets, but not used too frequently in deep learning

k-Nearest Neighbor on images **never** used.

- Very slow at test time
- Distance metrics on pixels are not informative
- Curse of dimensionality

Original



Boxed



Shifted



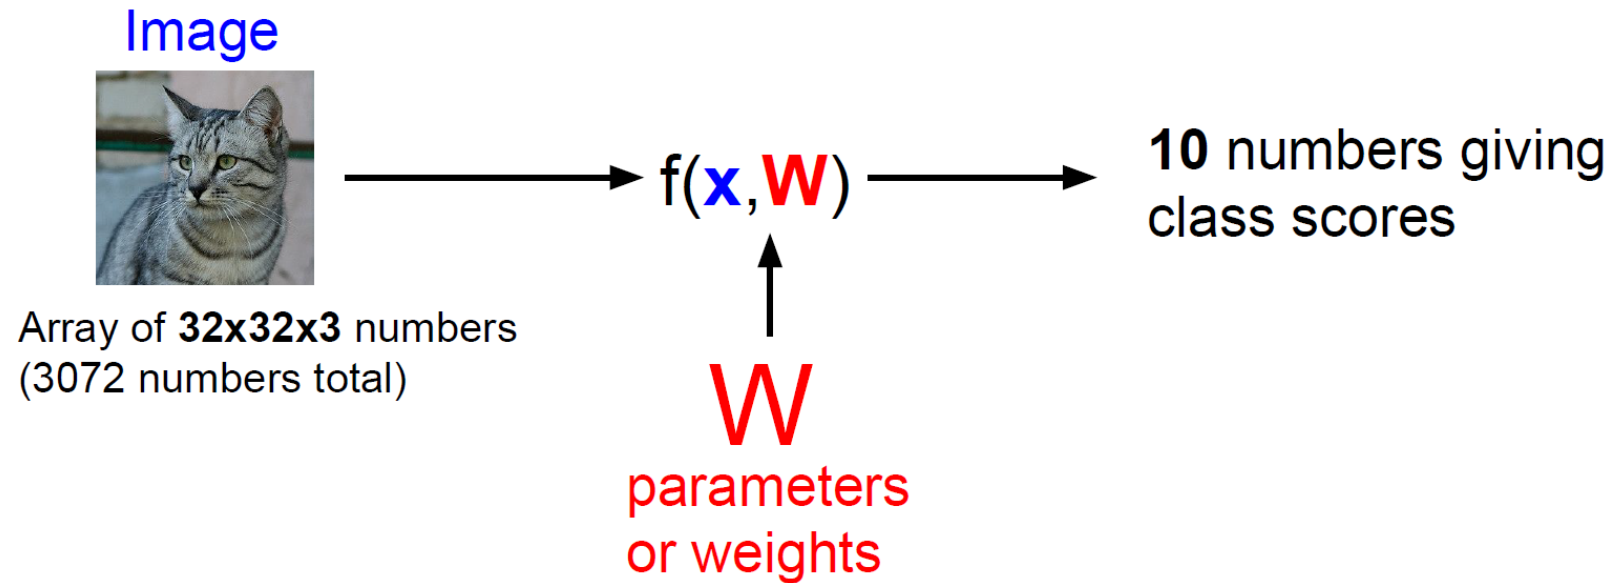
Tinted



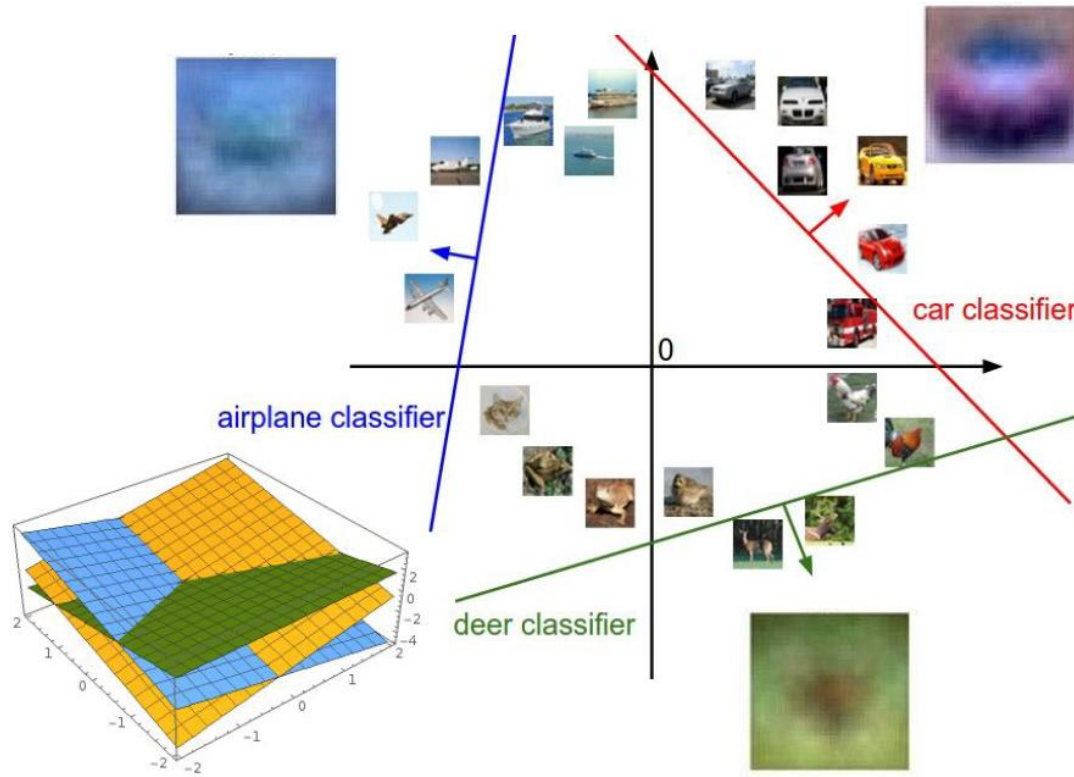
Linear Classification

- Components of neural networks

$$f(x, W) = Wx + b(ias) \quad f: 10by1, W: 10by3072, x: 3072by1, b: 10by1$$



Linear Classification



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers
(3072 numbers total)


Hard Case for a linear classification

Class 1:

number of pixels > 0 odd

Class 2:

number of pixels > 0 even




Class 1:

$1 \leq \text{L2 norm} \leq 2$

Class 2:

Everything else




Class 1:

Three modes

Class 2:

Everything else



Define $f(x, W)$

- Different performance to different class -> Is this good or bad?



airplane	-3.45	-0.51	3.42
automobile	-8.87	6.04	4.64
bird	0.09	5.31	2.65
cat	2.9	-4.22	5.1
deer	4.48	-4.19	2.64
dog	8.02	3.58	5.55
frog	3.78	4.49	-4.34
horse	1.06	-4.37	-1.5
ship	-0.36	-2.09	-4.79
truck	-0.72	-2.93	6.14

Thank you